

实验 2：内存管理

邬天行 522030910206

练习题 1: 完成 `kernel/mm/buddy.c` 中的 `split_chunk`、`merge_chunk`、`buddy_get_pages`、和 `buddy_free_pages` 函数中的 LAB 2 TODO 1 部分,其中 `buddy_get_pages` 用于分配指定阶大小的连续物理页,`buddy_free_pages` 用于释放已分配的连续物理页。

考虑 `buddy.h` 和 `list.h` 两个文件中关于buddy system相关的代码:

```
struct page {
    /* Free list */
    // 该物理页在链表中的指代节点
    struct list_head node;
    /* Whether the correspond physical page is free now. */
    // 标记是否被分配。 0表示未分配,1表示已分配
    int allocated;
    /* The order of the memory chunk that this page belongs to. */
    // 该物理页的阶次,也即大小为 2^order * PAGE_SIZE
    int order;
    /* Used for ChCore slab allocator. */
    void *slab;
    /* The physical memory pool this page belongs to */
    struct phys_mem_pool *pool;
};

struct free_list {
    //链表的表头,链表内每一项均指向一个未被分配的物理页
    struct list_head free_list;
    //该链表的长度,也即可分配的物理页个数
    unsigned long nr_free;
};

//一个双向链表结构,其中表头节点不行使功能
struct list_head {
    struct list_head *prev;
    struct list_head *next;
};

//链表相关的函数
static inline void init_list_head(struct list_head *list);
static inline void list_add(struct list_head *new, struct list_head *head);
static inline void list_append(struct list_head *new, struct list_head *head);
static inline void list_del(struct list_head *node);
static inline bool list_empty(struct list_head *head);

/*
list_entry(ptr,type,field)返回:
ptr作为[type结构体的field成员]的指针 对应的 type结构体指针
*/
```

具体来说,ptr为指向结构体某个成员的指针, type为该结构体类型, field为这个成员
原理为利用结构体分配内存连续性, container_of宏先通过((type *) (0))->field获得field的内存offset, 再将与field对应的ptr减去这个offset从而获得type本身的指针
*/

```
#define list_entry(ptr, type, field) \
    container_of(ptr, type, field)

#define container_of(ptr, type, field) \
    ((type *) ((void *) (ptr) - (void *) (&(((type *) (0))->field))))
```

涉及页面分裂和合并的函数维护free_lists链表均注意需要同时维护链表本体和长度, 且需要维护allocated和order成员

对于 split_chunk 函数: 采用递归方式实现

递归终止条件: 当前页面的阶次已经是我们需要的阶次

对于每一个chunk, 获取比它阶次低一级的chunk的地址作为buddy_chunk送入free_lists, 然后将分裂后的chunk继续递归

```
__maybe_unused static struct page *split_chunk(struct phys_mem_pool
*__maybe_unused pool,
                                                int __maybe_unused order,
                                                struct page *__maybe_unused chunk)
{
    /* LAB 2 TODO 1 BEGIN */
    /*
     * Hint: Recursively put the buddy of current chunk into
     * a suitable free list.
     */
    /* BLANK BEGIN */

    // end of recursion
    if (chunk->order == order) return chunk;

    chunk->order--;
    struct page *buddy_chunk = get_buddy_chunk(pool, chunk);
    buddy_chunk->order = chunk->order;
    buddy_chunk->allocated = 0;
    list_add(&(buddy_chunk->node), &(pool->free_lists[buddy_chunk->order].free_list));
    pool->free_lists[buddy_chunk->order].nr_free++;
    return split_chunk(pool, order, chunk);
    /* BLANK END */
    /* LAB 2 TODO 1 END */
}
```

对于 merge_chunk 函数: 采用递归方式实现

递归终止条件: 已经合并到最大阶次

注意到`get_buddy_chunk`函数仅能返回单一地址, 需要检查如下情况:

- 该页面是否存在
- 是否与原chunk具有相同的阶次
- 是否处于未分配状态

对于每一个chunk, 将`buddy_chunk`从`free_lists`内移除, 并取二者更低的地址作为高一阶次的chunk的地址继续递归合并, 并维护相应信息

```
__maybe_unused static struct page * merge_chunk(struct phys_mem_pool
* __maybe_unused pool,
                                         struct page * __maybe_unused chunk)
{
    /* LAB 2 TODO 1 BEGIN */
    /*
     * Hint: Recursively merge current chunk with its buddy
     * if possible.
     */
    /* BLANK BEGIN */

    // end of recursion
    if (chunk->order == (BUDDY_MAX_ORDER-1)) return chunk;

    //get buddy_chunk
    struct page *buddy_chunk = get_buddy_chunk(pool, chunk);

    //false conditions:
    //1. no else buddy chunk
    if (!buddy_chunk) return chunk;
    //2. buddy chunk can't be allocated
    if (buddy_chunk->allocated == 1) return chunk;
    //3. buddy chunk is not free as a whole
    if (buddy_chunk->order != chunk->order) return chunk;

    list_del(&(buddy_chunk->node));
    pool->free_lists[buddy_chunk->order].nr_free--;

    buddy_chunk->order++;
    chunk->order++;

    //let chunk be the addr of smaller one
    if (chunk > buddy_chunk) chunk = buddy_chunk;
    return merge_chunk(pool, chunk);

    /* BLANK END */
    /* LAB 2 TODO 1 END */
}
```

对于 `buddy_get_pages` 函数, 我们从当前阶次向上遍历, 找到阶次足够大的chunk后再分裂至需要的阶次, 并维护相应信息

注意如果找不到对应页面需要goto out以防止进程锁未释放

```

struct page *buddy_get_pages(struct phys_mem_pool *pool, int order)
{
    int cur_order;
    struct list_head *free_list;
    struct page *page = NULL;

    if (unlikely(order >= BUDDY_MAX_ORDER)) {
        kwarn("ChCore does not support allocating such too large "
            "contiguous physical memory\n");
        return NULL;
    }

    lock(&pool->buddy_lock);

    /* LAB 2 TODO 1 BEGIN */
    /*
     * Hint: Find a chunk that satisfies the order requirement
     * in the free lists, then split it if necessary.
     */
    /* BLANK BEGIN */
    for(cur_order = order; cur_order < BUDDY_MAX_ORDER; cur_order++)
    {
        if (pool->free_lists[cur_order].nr_free > 0)
        {
            free_list = &pool->free_lists[cur_order].free_list;
            page = list_entry(free_list->next, struct page, node);
            break;
        }
    }

    if (!page) goto out;

    list_del(&page->node);
    pool->free_lists[cur_order].nr_free--;
    page = split_chunk(pool, order, page);
    page->allocated = 1;

    /* BLANK END */
    /* LAB 2 TODO 1 END */
out: __maybe_unused
    unlock(&pool->buddy_lock);
    return page;
}

```

对于 `buddy_free_pages` 函数,我们直接合并,并维护相应信息

```

void buddy_free_pages(struct phys_mem_pool *pool, struct page *page)
{

```

```

int order;
struct list_head *free_list;
lock(&pool->buddy_lock);

/* LAB 2 TODO 1 BEGIN */
/*
 * Hint: Merge the chunk with its buddy and put it into
 * a suitable free list.
 */
/* BLANK BEGIN */
page->allocated = 0;
page = merge_chunk(pool, page);
free_list = &(pool->free_lists[page->order].free_list);
list_add(&(page->node), free_list);
pool->free_lists[page->order].nr_free++;
/* BLANK END */
/* LAB 2 TODO 1 END */

unlock(&pool->buddy_lock);
}

```

练习题 2: 完成 `kernel/mm/slab.c` 中的 `choose_new_current_slab`、`alloc_in_slab_impl` 和 `free_in_slab` 函数中的 LAB 2 TODO 2 部分,其中 `alloc_in_slab_impl` 用于在 slab 分配器中分配指定阶大小的内存,而 `free_in_slab` 则用于释放上述已分配的内存。

考虑 `slab.h` 文件中关于 slab 的代码:

```

struct slab_header {
    /* The list of free slots, which can be converted to struct
    slab_slot_list. */
    //当前slab下slab_slot_list链表的表头,与前面的双向链表的区别在于该表头可以正常行
    驶功能,且可以被强制类型转换为slab_slot_list类型
    void *free_list_head;
    /* Partial slab list. */
    //指向下一个slab_slot_list链表的表头
    struct list_head node;

    //slab对应的page的阶次
    int order;
    //最大自由slot数
    unsigned short total_free_cnt; /* MAX: 65536 */
    //当前自由slot数,关联free_list_head链表的长度
    unsigned short current_free_cnt;
};

/* Each free slot in one slab is regarded as slab_slot_list. */
struct slab_slot_list {
    //next_free指向当前slot_list下一个空的slot
    void *next_free;
};

```

```

struct slab_pointer {
    //当前使用的slab_slot_list
    struct slab_header *current_slab;
    //所有partial_slab_list的开头组成前文的双向链表
    struct list_head partial_slab_list;
};

```

对于 `choose_new_current_slab` 函数,我们查找`partial_slab_list`,从中直接取第一项作为新的`current_slab`即可,否则返回NULL

```

static void choose_new_current_slab(struct slab_pointer * __maybe_unused pool)
{
    /* LAB 2 TODO 2 BEGIN */
    /* Hint: Choose a partial slab to be a new current slab. */
    /* BLANK BEGIN */
    struct list_head *list = &(amp;pool->partial_slab_list);
    if (list_empty(list))
    {
        pool->current_slab = NULL;
    }
    else
    {
        pool->current_slab = (struct slab_header*)list_entry(list->next,
struct slab_header, node);
        list_del(list->next);
    }
    /* BLANK END */
    /* LAB 2 TODO 2 END */
}

```

对于 `alloc_in_slab_impl` 函数,我们从`current_slab`中取出第一个slot作为分配的内存,并维护相关变量,若此时`current_slab`的slot均已用尽,则调用 `choose_new_current_slab` 函数重新分一个`current_slab`

```

static void *alloc_in_slab_impl(int order)
{
    .....

    /* LAB 2 TODO 2 BEGIN */
    /*
     * Hint: Find a free slot from the free list of current slab.
     * If current slab is full, choose a new slab as the current one.
     */
    /* BLANK BEGIN */
    free_list = (struct slab_slot_list *) (current_slab->free_list_head);
    next_slot = free_list->next_free;
    current_slab->free_list_head = next_slot;
    current_slab->current_free_cnt--;
    if (current_slab->current_free_cnt == 0)

```

```

choose_new_current_slab(&slab_pool[order]);
    /* BLANK END */
    /* LAB 2 TODO 2 END */

    .....

}

```

对于 `free_in_slab` 函数,我们将待释放的slot链表节点重新加入`free_list_head`链表中,放在表头最为简单,并维护长度

```

void free_in_slab(void *addr)
{
    .....

    /* LAB 2 TODO 2 BEGIN */
    /*
     * Hint: Free an allocated slot and put it back to the free list.
     */
    /* BLANK BEGIN */
    slot->next_free = slab->free_list_head;
    slab->free_list_head = (void *)slot;
    slab->current_free_cnt++;
    /* BLANK END */
    /* LAB 2 TODO 2 END */

    .....

}

```

练习题 3: 完成 `kernel/mm/kmalloc.c` 中的 `_kmalloc` 函数中的 LAB 2 TODO 3 部分,在适当位置调用对应的函数,实现 `kmalloc` 功能

以`SLAB_MAX_SIZE`作为分界线:

- 当申请内存的大小不超过一个page/slab的大小,调用 `alloc_in_slab` 分配小块内存
- 否则调用 `get_pages` 函数从buddy system中分配大块内存

```

void *_kmalloc(size_t size, bool is_record, size_t *real_size)
{
    .....

    if (size <= SLAB_MAX_SIZE) {
        /* LAB 2 TODO 3 BEGIN */
        /* Step 1: Allocate in slab for small requests. */
        /* BLANK BEGIN */
        addr = alloc_in_slab(size, real_size);
        /* BLANK END */
    }
    #if ENABLE_MEMORY_USAGE_COLLECTING == ON
        if(is_record && collecting_switch) {

```

```

        record_mem_usage(*real_size, addr);
    }
#endif
} else {
    /* Step 2: Allocate in buddy for large requests. */
    /* BLANK BEGIN */
    order = size_to_page_order(size);
    addr = _get_pages(order, is_record);
    /* BLANK END */
    /* LAB 2 TODO 3 END */
}

.....
}

```

练习题 4: 完成 `kernel/arch/aarch64/mm/page_table.c` 中的 `query_in_pgtbl`、`map_range_in_pgtbl_common`、`unmap_range_in_pgtbl` 和 `mprotect_in_pgtbl` 函数中的 LAB 2 TODO 4 部分,分别实现页表查询、映射、取消映射和修改页表权限的操作,以 4KB 页为粒度。

考虑接下来均需要使用的 `get_next_ptp` 函数:

```

static int get_next_ptp(ptp_t *cur_ptp, u32 level, vaddr_t va, ptp_t **next_ptp,
                        pte_t **pte, bool alloc, __maybe_unused long *rss)
{
    u32 index = 0;
    pte_t *entry;

    if (cur_ptp == NULL)
        return -ENOMAPPING;

    switch (level) {
    case L0:
        index = GET_L0_INDEX(va);
        break;
    case L1:
        index = GET_L1_INDEX(va);
        break;
    case L2:
        index = GET_L2_INDEX(va);
        break;
    case L3:
        index = GET_L3_INDEX(va);
        break;
    default:
        BUG("unexpected level\n");
        return -EINVAL;
    }

    entry = &(cur_ptp->ent[index]);
}

```



```

    if (IS_PTE_INVALID(entry->pte)) {
        if (alloc == false) {
            return -ENOMAPPING;
        } else {
            /* alloc a new page table page */
            ptp_t *new_ptp;
            paddr_t new_ptp_paddr;
            pte_t new_pte_val;

            /* alloc a single physical page as a new page table page
             */
            new_ptp = get_pages(0);
            if (new_ptp == NULL)
                return -ENOMEM;
            if (rss) *rss += PAGE_SIZE;
            memset((void *)new_ptp, 0, PAGE_SIZE);

            new_ptp_paddr = virt_to_phys((vaddr_t)new_ptp);

            new_pte_val.pte = 0;
            new_pte_val.table.is_valid = 1;
            new_pte_val.table.is_table = 1;
            new_pte_val.table.next_table_addr = new_ptp_paddr
                                                >> PAGE_SHIFT;

            /* same effect as: cur_ptp->ent[index] = new_pte_val; */
            entry->pte = new_pte_val.pte;
        }
    }

    *next_ptp = (ptp_t *)GET_NEXT_PTP(entry);
    *pte = entry;
    if (IS_PTE_TABLE(entry->pte))
        return NORMAL_PTP;
    else
        return BLOCK_PTP;
}

```

该函数有几种不同的返回值:

- -ENOMAPPING, 表明在查询时没找到下一级页表项/大页
- -ENOMEM, 表明物理内存已用完,本实验中暂不需考虑
- NORMAL_PTP, 此时找到的是下一级页表项
- BLOCK_PTP,此时找到的是是下一级大页

对于 `query_in_pgtbl` 函数,我们不断调用 `get_next_ptp` 函数,直到找到相应的页表项/大页,接着利用提供的相关宏通过分开计算pfn和offset得到具体的物理地址

```

int query_in_pgtbl(void *pgtbl, vaddr_t va, paddr_t *pa, pte_t **entry)
{
    /* LAB 2 TODO 4 BEGIN */

```

```

/*
 * Hint: Walk through each level of page table using `get_next_ptp`,
 * return the pa and pte until a L2/L3 block or page, return
 * `-ENOMAPPING` if the va is not mapped.
 */
/* BLANK BEGIN */
ptp_t *l0_ptp = (ptp_t *) pgtbl;
ptp_t *l1_ptp = NULL, *l2_ptp = NULL, *l3_ptp = NULL;
ptp_t *phys_page;
pte_t *pte;
int ret;

ret = get_next_ptp(l0_ptp, L0, va, &l1_ptp, &pte, 0, NULL);
if (ret < 0) return ret;

ret = get_next_ptp(l1_ptp, L1, va, &l2_ptp, &pte, 0, NULL);
if (ret < 0)
{
    return ret;
}
else if (ret == BLOCK_PTP)
{
    *pa = (pte->l1_block.pfn << L1_INDEX_SHIFT) |
(GET_VA_OFFSET_L1(va));
    if (entry) *entry = pte;
    return 0;
}
ret = get_next_ptp(l2_ptp, L2, va, &l3_ptp, &pte, 0, NULL);
if (ret < 0)
{
    return ret;
}
else if (ret == BLOCK_PTP)
{
    *pa = (pte->l2_block.pfn << L2_INDEX_SHIFT) |
(GET_VA_OFFSET_L2(va));
    if (entry) *entry = pte;
    return 0;
}

ret = get_next_ptp(l3_ptp, L3, va, &phys_page, &pte, 0, NULL);
if (ret < 0) return ret;

*pa = (pte->l3_page.pfn << L3_INDEX_SHIFT) | (GET_VA_OFFSET_L3(va));
if (entry) *entry = pte;

/* BLANK END */
/* LAB 2 TODO 4 END */
return 0;
}

```

对于 `map_range_in_pgtbl_common`, `unmap_range_in_pgtbl` 和 `mprotect_in_pgtbl` 三个函数, 逻辑均为通过 `get_next_ptp` 函数找到需要映射/解除映射/设置pte的物理页, 然后执行对应操作; 注意前两个函数需要在分配/解除物理页时计算对应的rss增减量

```
static int map_range_in_pgtbl_common(void *pgtbl, vaddr_t va, paddr_t pa,
                                     size_t len, vmr_prop_t flags, int kind,
                                     __maybe_unused long *rss)
{
    /* LAB 2 TODO 4 BEGIN */
    /*
     * Hint: Walk through each level of page table using `get_next_ptp`,
     * create new page table page if necessary, fill in the final level
     * pte with the help of `set_pte_flags`. Iterate until all pages are
     * mapped.
     * Since we are adding new mappings, there is no need to flush TLBs.
     * Return 0 on success.
     */
    /* BLANK BEGIN */
    s64 total_pg_cnt = DIV_ROUND_UP(len, PAGE_SIZE);
    ptp_t *l0_ptp = (ptp_t *) pgtbl;
    ptp_t *l1_ptp = NULL, *l2_ptp = NULL, *l3_ptp = NULL;
    pte_t *pte;
    int ret;

    while (total_pg_cnt > 0)
    {
        ret = get_next_ptp(l0_ptp, L0, va, &l1_ptp, &pte, 1, rss);
        if (ret) return ret;
        ret = get_next_ptp(l1_ptp, L1, va, &l2_ptp, &pte, 1, rss);
        if (ret)
        {
            return ret;
        }
        else if (ret == BLOCK_PTP)
        {
            total_pg_cnt -= L1_PER_ENTRY_PAGES;
            if (rss) *rss += L1_PER_ENTRY_PAGES * PAGE_SIZE;
            continue;
        }

        ret = get_next_ptp(l2_ptp, L2, va, &l3_ptp, &pte, 1, rss);
        if (ret)
        {
            return ret;
        }
        else if (ret == BLOCK_PTP)
        {
            total_pg_cnt -= L2_PER_ENTRY_PAGES;
            if (rss) *rss += L2_PER_ENTRY_PAGES * PAGE_SIZE;
            continue;
        }
    }
}
```

```

        total_pg_cnt -= L3_PER_ENTRY_PAGES;
        pte = &(l3_ptp->ent[GET_L3_INDEX(va)]);
        set_pte_flags(pte, flags, kind);
        pte->l3_page.is_page = 1;
        pte->l3_page.is_valid = 1;
        pte->l3_page.pfn = pa >> L3_INDEX_SHIFT;

        va += PAGE_SIZE;
        pa += PAGE_SIZE;
        if (rss) *rss += L3_PER_ENTRY_PAGES * PAGE_SIZE;
    }
    /* BLANK END */
    /* LAB 2 TODO 4 END */
    dsb(ishst);
    isb();
    return 0;
}

```

```

int unmap_range_in_pgtbl(void *pgtbl, vaddr_t va, size_t len,
                        __maybe_unused long *rss)
{
    /* LAB 2 TODO 4 BEGIN */
    /*
     * Hint: Walk through each level of page table using `get_next_ptp`,
     * mark the final level pte as invalid. Iterate until all pages are
     * unmapped.
     * You don't need to flush tlb here since tlb is now flushed after
     * this function is called.
     * Return 0 on success.
     */
    /* BLANK BEGIN */

    s64 total_pg_cnt = DIV_ROUND_UP(len, PAGE_SIZE);
    ptp_t *l0_ptp = (ptp_t *) pgtbl;
    ptp_t *l1_ptp = NULL, *l2_ptp = NULL, *l3_ptp = NULL, *phys_page = NULL;
    pte_t *pte;
    int ret;

    while (total_pg_cnt)
    {
        ret = get_next_ptp(l0_ptp, L0, va, &l1_ptp, &pte, 0, rss);
        if (ret < 0) return ret;
        ret = get_next_ptp(l1_ptp, L1, va, &l2_ptp, &pte, 0, rss);
        if (ret < 0)
        {
            return ret;
        }
        else if (ret == BLOCK_PTP)
        {
            total_pg_cnt -= L1_PER_ENTRY_PAGES;
            pte->pte = PTE_DESCRIPTOR_INVALID;
        }
    }
}

```

```

        continue;
    }
    ret = get_next_ptp(l2_ptp, L2, va, &l3_ptp, &pte, 0, rss);
    if (ret < 0)
    {
        return ret;
    }
    else if (ret == BLOCK_PTP)
    {
        total_pg_cnt -= L2_PER_ENTRY_PAGES;
        pte->pte = PTE_DESCRIPTOR_INVALID;
        continue;
    }
    ret = get_next_ptp(l3_ptp, L3, va, &phys_page, &pte, 0, rss);
    if (ret < 0) return ret;

    pte->pte = PTE_DESCRIPTOR_INVALID;
    if(rss) *rss -= PAGE_SIZE;
    recycle_pgtable_entry(l0_ptp, l1_ptp, l2_ptp, l3_ptp, va, rss);
    va += PAGE_SIZE;
    total_pg_cnt--;
    if (total_pg_cnt == 0) break;
}
/* BLANK END */
/* LAB 2 TODO 4 END */

dsb(ishst);
isb();

return 0;
}

```

```

int mprotect_in_pgtbl(void *pgtbl, vaddr_t va, size_t len, vmr_prop_t flags)
{
    /* LAB 2 TODO 4 BEGIN */
    /*
     * Hint: Walk through each level of page table using `get_next_ptp`,
     * modify the permission in the final level pte using `set_pte_flags`.
     * The `kind` argument of `set_pte_flags` should always be `USER_PTE`.
     * Return 0 on success.
     */
    /* BLANK BEGIN */
    s64 total_pg_cnt = DIV_ROUND_UP(len, PAGE_SIZE);
    ptp_t *l0_ptp = (ptp_t *) pgtbl;
    ptp_t *l1_ptp = NULL, *l2_ptp = NULL, *l3_ptp = NULL, *phys_page = NULL;
    pte_t *pte;
    int ret;

    while (total_pg_cnt)
    {
        ret = get_next_ptp(l0_ptp, L0, va, &l1_ptp, &pte, 0, NULL);
    }
}

```

```

        if (ret < 0) return ret;
        ret = get_next_ptp(l1_ptp, L1, va, &l2_ptp, &pte, 0, NULL);
        if (ret < 0)
        {
            return ret;
        }
        else if (ret == BLOCK_PTP)
        {
            total_pg_cnt -= L1_PER_ENTRY_PAGES;
            continue;
        }
        ret = get_next_ptp(l2_ptp, L2, va, &l3_ptp, &pte, 0, NULL);
        if (ret < 0)
        {
            return ret;
        }
        else if (ret == BLOCK_PTP)
        {
            total_pg_cnt -= L2_PER_ENTRY_PAGES;
            continue;
        }
        ret = get_next_ptp(l3_ptp, L3, va, &phys_page, &pte, 0, NULL);
        if (ret < 0) return ret;

        set_pte_flags(pte, flags, USER_PTE);
        va += PAGE_SIZE;
        total_pg_cnt--;
        if (total_pg_cnt == 0) break;
    }
    /* BLANK END */
    /* LAB 2 TODO 4 END */
    return 0;
}

```

思考题 5: 阅读 Arm Architecture Reference Manual,思考要在操作系统中支持写时拷贝(Copy-on-Write,CoW)需要配置页表描述符的哪个/哪些字段,并在发生页错误时如何处理。(在完成第三部分后,你也可以阅读页错误处理的相关代码,观察 ChCore 是如何支持 Cow 的)

为了支持写时拷贝,需要在页表项中将 AP(Access Permissions)字段配置为只读。

考虑pgfault_handler.c中处理CoW的代码:

```

/*
 * Perform general COW
 * Step-1: get PA of page containing fault_addr, so as kernel VA of that page
 * Step-2: allocate a new page and record in VMR
 * Step-3: copy using kernel VA to new page
 * Step-4(?): update VMR perm (How and when? Neccessary?)
 * Step-5: update PTE permission and PPN
 * Step-6: Flush TLB of user virtual page(user_vpa)

```

```

*/
static int __do_general_cow(struct vmSPACE *vmSPACE, struct vmregion *vmr,
                           vaddr_t fault_addr, pte_t *fault_pte,
                           struct common_pte_t *pte_info)
{
    vaddr_t kva, user_vpa;
    void *new_page;
    paddr_t new_pa;
    struct common_pte_t new_pte_attr;
    int ret = 0;

    /* Step-1: get PA of page containing fault_addr, so as kernel VA of that
     * page */
    kva = phys_to_virt(pte_info->ppn << PAGE_SHIFT);

    /* Step-2: allocate a new page and record in VMR */
    new_page = get_pages(0);
    if (!new_page) {
        ret = -ENOMEM;
        goto out;
    }
    new_pa = virt_to_phys(new_page);

    ret = vmregion_record_cow_private_page(vmr, fault_addr, new_page);
    if (ret)
        goto out_free_page;

    /* Step-3: copy using kernel VA to new page */
    memcpy(new_page, (void *)kva, PAGE_SIZE);

    /* Step-5: update PTE permission and PPN */
    new_pte_attr.ppn = new_pa >> PAGE_SHIFT;
    new_pte_attr.perm = pte_info->perm | VMR_WRITE;
    new_pte_attr.valid = 1;
    new_pte_attr.access = 0;
    new_pte_attr.dirty = 0;

    update_pte(fault_pte, L3, &new_pte_attr);

    /* Step-6: Flush TLB of user virtual page(user_vpa) */
    user_vpa = ROUND_DOWN(fault_addr, PAGE_SIZE);
    flush_tlb_by_range(vmSPACE, user_vpa, PAGE_SIZE);

    return 0;
out_free_page:
    free_pages(new_page);
out:
    return ret;
}

```

当只读页内发生写操作时,大致经历这些流程:

- 触发缺页异常(在调用相关函数时进入__do_general_cow函数)

- 取得发生异常的物理地址,并分配一个新的物理页
- 将共享页的内容拷贝至新的物理页中
- 将这个物理页的访问权限配置为可读可写
- 刷新TLB

思考题 6: 为了简单起见,在 ChCore 实验 Lab1 中没有为内核页表使用细粒度的映射,而是直接沿用了启动时的粗粒度页表,请思考这样做有什么问题。

- 为内核分配的大页通常不会被完全使用,从而产生大量内部碎片,造成内存浪费
- 需要访问小块数据时,可能需要映射和访问多个不相关的内存块
- 缺少详细的权限控制

挑战题 7: 使用前面实现的 `page_table.c` 中的函数,在内核启动后的 `main` 函数中重新配置内核页表,进行细粒度的映射。

从lab1中我们可以看到对应的映射关系(不考虑映射粒度):

因此,在内核启动时,首先需要对外核自身、其余可用物理内存和外设内存进行虚拟地址映射,最简单的映射方式是一对一的映射,即将虚拟地址 `0xffff_0000_0000_0000 + addr` 映射到 `addr`。需要注意的是,在 ChCore 实验中我们使用了 `0xffff_ff00_0000_0000` 作为内核虚拟地址的开始(注意开头 `f` 数量的区别),不过这不影响我们对知识点的理解。

在树莓派 3B+ 机器上,物理地址空间分布如下¹⁰:

¹⁰ [bcm2836-peripherals.pdf](#) & [Raspberry Pi Hardware - Peripheral Addresses](#)

物理地址范围	对应设备
0x00000000 ~ 0x3f000000	物理内存 (SDRAM)
0x3f000000 ~ 0x40000000	共享外设内存
0x40000000 ~ 0xffffffff	本地 (每个 CPU 核独立) 外设内存

现在将目光转移到 `kernel/arch/aarch64/boot/raspi3/init/mmu.c` 文件,我们需要在 `init_kernel_pt` 为内核配置从 `0x00000000` 到 `0x80000000` (`0x40000000` 后的 1G, ChCore 只需使用这部分地址中的本地外设) 的映射,其中 `0x00000000` 到 `0x3f000000` 映射为 `normal memory`, `0x3f000000` 到 `0x80000000` 映射为 `device memory`,其中 `0x00000000` 到 `0x40000000` 以 2MB 块粒度映射, `0x40000000` 到 `0x80000000` 以 1GB 块粒度映射。

在mmu启动后需要经历如下流程:

- 先重新映射kernel
- 申请一块页表,将基地址作为ttbr_el1的值
- 按照对应的物理地址范围分别映射物理内存和设备内存
- 刷新TLB

```
void main(paddr_t boot_flag, void *info)
{
    .....
}
```



```

/* Mapping KSTACK into kernel page table. */
map_range_in_pgtbl_kernel((void*)((unsigned long)boot_ttbr1_l0 + KBASE),
    KSTACKx_ADDR(0),
    (unsigned long)(cpu_stacks[0]) - KBASE,
    CPU_STACK_SIZE, VMR_READ | VMR_WRITE);

#define PHYSICAL_START (0x0UL)
#define PERIPHERAL_BASE (0x3F000000UL)
#define PHYSICAL_END (0xffffffffUL)
#define KERNEL_BASE (0xffffffff00000000UL)
void *ttbr1_el1 = get_pages(0);
//map for kernel
map_range_in_pgtbl_kernel(ttbr1_el1,
    KSTACKx_ADDR(0),
    virt_to_phys(cpu_stacks[0]),
    CPU_STACK_SIZE, VMR_READ | VMR_WRITE);

//map for SDRAM
map_range_in_pgtbl_kernel(ttbr1_el1,
    KERNEL_BASE + PHYSICAL_START,
    PHYSICAL_START,
    PERIPHERAL_BASE - PHYSICAL_START, VMR_EXEC);
//map for peripheral memory
map_range_in_pgtbl_kernel(ttbr1_el1,
    KERNEL_BASE + PERIPHERAL_BASE,
    PERIPHERAL_BASE,
    PHYSICAL_END - PERIPHERAL_BASE, VMR_DEVICE);
flush_tlb_all();
kinfo("[ChCore] kernel remap finished\n");
#undef PHYSICAL_START
#undef PERIPHERAL_BASE
#undef PHYSICAL_END
#undef KERNEL_BASE

/* Init exception vector */
arch_interrupt_init();
timer_init();
kinfo("[ChCore] interrupt init finished\n");

.....
}

```

练习题 8: 完成 `kernel/arch/aarch64/irq/pgfault.c` 中的 `do_page_fault` 函数中的 LAB 2 TODO 5 部分,将缺页异常转发给 `handle_trans_fault` 函数。

直接将 `handle_trans_fault` 函数需要的参数传入即可

```

void do_page_fault(u64 esr, u64 fault_ins_addr, int type, u64 *fix_addr)
{
    .....
    fault_addr = get_fault_addr();
}

```

```

fsc = GET_ESR_EL1_FSC(esr);
switch (fsc) {
case DFSC_TRANS_FAULT_L0:
case DFSC_TRANS_FAULT_L1:
case DFSC_TRANS_FAULT_L2:
case DFSC_TRANS_FAULT_L3: {
    /* LAB 2 TODO 5 BEGIN */
    /* BLANK BEGIN */
    ret = handle_trans_fault(current_thread->vmSPACE, fault_addr);
    /* BLANK END */
    /* LAB 2 TODO 5 END */
    .....
}
.....
}
.....
}

```

练习题 9: 完成 `kernel/mm/vmSPACE.c` 中的 `find_vmr_for_va` 函数中的 LAB 2 TODO 6 部分, 找到一个虚拟地址找在其虚拟地址空间中的 VMR。

利用 `rb_search` 函数查找虚拟地址所在的树上节点, 然后利用 `rb_entry` 宏取出相应的 `vmr` 即可

```

__maybe_unused struct vmregion *find_vmr_for_va(struct vmSPACE *vmSPACE,
                                                  vaddr_t addr)
{
    /* LAB 2 TODO 6 BEGIN */
    /* Hint: Find the corresponding vmr for @addr in @vmSPACE */
    /* BLANK BEGIN */
    struct rb_node *addr_node = rb_search(&vmSPACE->vmr_tree, (const void
*)addr, cmp_vmr_and_va);
    struct vmregion *vmr = NULL;
    if (addr_node) vmr = rb_entry(addr_node, struct vmregion, tree_node);
    return vmr;
    /* BLANK END */
    /* LAB 2 TODO 6 END */
}

```

练习题 10: 完成 `kernel/mm/pgfault_handler.c` 中的 `handle_trans_fault` 函数中的 LAB 2 TODO 7 部分(函数内共有 3 处填空, 不要遗漏), 实现 `PMO_SHM` 和 `PMO_ANONYM` 的按需物理页分配。你可以阅读代码注释, 调用你之前见到过的相关函数来实现功能。

- 若物理页尚未分配, 则调用 `get_pages` 函数申请一个物理页, 然后将其内容清零并添加映射
- 否则修改相应的页表映射即可

```

int handle_trans_fault(struct vmSPACE *vmSPACE, vaddr_t fault_addr)
{

```

```

.....
switch (pmo->type) {
case PMO_ANONYM:
case PMO_SHM: {
    .....
    if (pa == 0) {
        /*
         * Not committed before. Then, allocate the physical
         * page.
         */
        /* LAB 2 TODO 7 BEGIN */
        /* BLANK BEGIN */
        /* Hint: Allocate a physical page and clear it to 0. */
        void *va = get_pages(0);
        pa = virt_to_phys(va);
        memset(va, 0, PAGE_SIZE);
        /* BLANK END */
        /*
         * Record the physical page in the radix tree:
         * the offset is used as index in the radix tree
         */
        kdebug("commit: index: %ld, 0x%lx\n", index, pa);
        commit_page_to_pmo(pmo, index, pa);

        /* Add mapping in the page table */
        lock(&vmospace->pgtbl_lock);
        /* BLANK BEGIN */
        ret = map_range_in_pgtbl(vmospace->pgtbl, fault_addr, pa,
PAGE_SIZE, perm, &(vmospace->rss));
        /* BLANK END */
        unlock(&vmospace->pgtbl_lock);
    } else {
        if (pmo->type == PMO_SHM || pmo->type == PMO_ANONYM) {
            /* Add mapping in the page table */
            lock(&vmospace->pgtbl_lock);
            /* BLANK BEGIN */
            ret = map_range_in_pgtbl(vmospace->pgtbl,
fault_addr, pa, PAGE_SIZE, perm, &(vmospace->rss));
            /* BLANK END */
            /* LAB 2 TODO 7 END */
            unlock(&vmospace->pgtbl_lock);
        }
    }
    .....
}
.....
}
}

```

挑战题 11: 我们在`map_range_in_pgtbl_common`、`unmap_range_in_pgtbl` 函数中预留了没有被使用过的参数`rss`用来统计map映射中实际的物理内存使用量¹, 你需要修改相关的代码来通过Compute

physical memory测试,不实现该挑战题并不影响其他部分功能的实现及测试。如果你想检测是否通过此部分测试,需要修改.config中CHCORE_KERNEL_PM_USAGE_TEST为ON

由于.config文件由Lab2/kernel/config.cmake和其他相关的文件在build时产生,因此实际需要修改的文件为Lab2/kernel/config.cmake

对于map_range_in_pgtbl_common和unmap_range_in_pgtbl的修改见练习题4

注意到get_next_ptp函数会分配物理页(见练习题4),而在解除映射时会调用try_release_ptp函数释放,因此这两个函数也需要加入对应的计算代码

```
static int try_release_ptp(ptp_t *high_ptp, ptp_t *low_ptp, int index,
                          __maybe_unused long *rss)
{
    int i;

    for (i = 0; i < PTP_ENTRIES; i++) {
        if (low_ptp->ent[i].pte != PTE_DESCRIPTOR_INVALID) {
            return 0;
        }
    }

    BUG_ON(index < 0 || index >= PTP_ENTRIES);
    high_ptp->ent[index].pte = PTE_DESCRIPTOR_INVALID;
    kfree(low_ptp);
    if (rss) *rss -= PAGE_SIZE;

    return 1;
}
```

至此,运行 make qemu 可以正常进入 shell,运行 make grade 可以通过所有测试。

```
[lwip] TCP/IP initialized.
[lwip] Add netif 0x5d6f57a6eed0
[lwip] register server value = 0
Network-CP-Daemon: running at localhost:4096
```



```
Welcome to ChCore shell!
$
```

```
make[1]: Leaving directory '/home/os/os/homework1'
Grading lab 2 ...(may take 10 seconds)
=====
Allocate & free order 0: 5
Allocate & free each order: 5
Allocate & free all orders: 5
Allocate & free all memory: 5
kmalloc: 10
Map & unmap one page: 10
Map & unmap multiple pages: 10
Map & unmap huge range: 20
Compute physical memory-1: 1
Compute physical memory-2: 1
Compute physical memory-3: 3
Page fault: 30
Score: 105/100
=====
```